Project Report

Learn and Give Back to Society

Machine Learning Approaches for Thyroid Disease Detection

A Comparative Study Using Logistic Regression, Decision Tree, and Random Forest Models

Maihtaab Sidhu

LAURIER
WILFRID LAURIER UNIVERSITY

Maihtaab, Sidhu

Table of contents

List of Figures

List of Tables

Preface

This report presents the results of a machine learning project conducted as part of the CP322 course. The objective of the project was to explore how machine learning techniques can be applied to medical datasets to predict thyroid disease. The project focuses on comparing several supervised learning models and evaluating their predictive performance using clinical data from the UCI Thyroid Disease Dataset. The analysis demonstrates how data preprocessing, model selection, and evaluation metrics can be combined to build an effective machine learning pipeline.

Abstract

This project investigates the use of machine learning models to predict hyperthyroid disease using the Thyroid Disease Dataset from the UCI Machine Learning Repository. The objective of the study is to determine whether patient medical attributes can be used to accurately classify the presence of hyperthyroid conditions.

The dataset was preprocessed by handling missing values, encoding categorical variables, and preparing features for machine learning models. After preprocessing, the data was divided into training and testing sets for model evaluation.

Three supervised learning models were implemented and compared: Logistic Regression, Decision Tree, and Random Forest. Model performance was evaluated using accuracy, precision, recall, and F1-score.

The results show that all models achieved high predictive accuracy, with Random Forest providing the best overall performance. These findings demonstrate that machine learning techniques can effectively analyze medical datasets and assist in predicting thyroid disease conditions.

Chapter 1

Introduction

Thyroid disease is one of the most common endocrine disorders and can significantly affect a patient's metabolism, energy levels, and overall health. Hyperthyroidism occurs when the thyroid gland produces excessive thyroid hormones, which can lead to symptoms such as weight loss, increased heart rate, anxiety, and fatigue. Early detection and diagnosis of thyroid disorders are important for effective treatment and prevention of long-term health complications.

With the increasing availability of medical datasets, machine learning techniques have become valuable tools for analyzing healthcare data and identifying patterns that may assist with disease prediction. Machine learning models can learn complex relationships between clinical features and disease outcomes, allowing them to support medical decision-making and improve diagnostic accuracy.

This project explores the application of machine learning models to predict hyperthyroid disease using the Thyroid Disease Dataset from the UCI Machine Learning Repository. Several classification algorithms are implemented and evaluated in order to determine which model performs best for this prediction task.

1.1 Overview

The goal of this project is to apply supervised machine learning techniques to classify whether a patient has a hyperthyroid condition based on a set of medical attributes. The dataset contains multiple clinical features related to thyroid hormone measurements, patient characteristics, and diagnostic indicators.

The project follows a typical machine learning workflow. First, the dataset is cleaned and preprocessed by handling missing values, converting data types, and encoding categorical variables. The processed data is then split into training and testing sets to allow model evaluation on unseen data.

Three classification models are implemented and compared: Logistic Regression, Decision Tree, and Random Forest. Each model is trained on the training dataset and evaluated using performance metrics such as accuracy, precision, recall, and F1-score. The results are then analyzed to determine which model provides the best predictive performance.

1.1.1 Existing System

Traditional approaches for diagnosing thyroid disorders rely on laboratory tests and clinical evaluation performed by healthcare professionals. Physicians interpret hormone levels and other clinical indicators to determine whether a patient has a thyroid disorder. While these methods are effective, analyzing large amounts of medical data can be time-consuming and may not always reveal complex patterns between variables.

Machine learning techniques provide an alternative approach by automatically learning relationships within data. By training predictive models on historical patient data, machine learning algorithms can identify patterns that help classify disease conditions. These models can serve as decision-support tools that assist healthcare professionals in identifying potential thyroid disease cases more efficiently.

1.2 Objectives of Project (Must be clearly, precisely defined and must be covered in the work)

The primary objective of this project is to develop and evaluate machine learning models for predicting hyperthyroid disease using clinical data.

The specific objectives of the project are:

- To analyze and preprocess the Thyroid Disease Dataset obtained from the UCI Machine Learning Repository.
- To implement multiple supervised machine learning models including Logistic Regression, Decision Tree, and Random Forest classifiers.
- To train and evaluate these models using appropriate performance metrics such as accuracy, precision, recall, and F1-score.
- To compare the performance of different models and determine which algorithm provides the most effective prediction of thyroid disease.
- To demonstrate how machine learning techniques can be applied to medical datasets to support disease detection and classification.

Chapter 2

Pre-Processing and Exploratory Data Analysis

Before applying machine learning models, it is necessary to prepare and analyze the dataset to ensure that it is suitable for model training. Data preprocessing and exploratory data analysis (EDA) are important steps in the machine learning pipeline because the quality of the data directly affects model performance.

In this project, the Thyroid Disease dataset was cleaned and processed to handle missing values, convert variables into appropriate formats, and prepare the features for classification models. Exploratory data analysis was also performed to understand the characteristics of the dataset and identify patterns that may influence model predictions.

2.1 Dataset Collection

The dataset used in this project is the Thyroid Disease Dataset obtained from the UCI Machine Learning Repository. This dataset contains clinical and laboratory measurements related to thyroid conditions. The data includes attributes such as patient demographics, laboratory test results, and diagnostic indicators associated with thyroid disorders.

The dataset contains several thousand patient records and multiple attributes that describe medical conditions and hormone measurements. The target variable indicates whether a patient has a hyperthyroid-related condition or a negative diagnosis. Because the dataset includes both numerical and categorical variables, preprocessing steps are required before applying machine learning algorithms.

2.2 Data Pre-processing

The raw dataset required several preprocessing steps before it could be used for training machine learning models.

First, missing values represented by the symbol "?" were replaced with null values. Handling missing values is important because many machine learning algorithms cannot process incomplete data. Numerical columns containing missing values were filled using the median of the respective feature, which helps reduce the influence of extreme values. For categorical variables, missing values were replaced with the most frequent category.

Next, categorical features were encoded using one-hot encoding so that they could be interpreted by machine learning algorithms. One-hot encoding converts categorical values into binary columns, allowing models to process non-numerical attributes effectively.

Data type conversion was also performed to ensure that numerical attributes were properly treated as numeric variables. This step ensures that statistical operations and model training processes work correctly.

After preprocessing, the dataset was divided into training and testing sets using an 80/20 split. The training dataset was used to train the machine learning models, while the testing dataset was used to evaluate the predictive performance of the models.

2.3 Exploratory Data Analysis and Visualisations

Exploratory data analysis was conducted to better understand the dataset and identify important characteristics that could influence model performance.

Initial analysis showed that the dataset contains a class imbalance, where negative thyroid cases significantly outnumber hyperthyroid cases. This imbalance can affect model performance, particularly for recall and precision metrics, because models may favor predicting the majority class.

Basic statistical analysis was used to examine the distribution of numerical variables and identify potential anomalies in the data. Understanding these distributions helps determine whether additional transformations or preprocessing steps are required.

Feature relationships were also explored to understand how different medical attributes may contribute to thyroid disease prediction. Machine learning models such as Decision Trees and Random Forests are particularly useful for identifying complex relationships between variables and disease outcomes.

Although advanced dimensionality reduction techniques were not required for this project, feature preprocessing and encoding ensured that the dataset was properly structured for model training and evaluation.

2.4 Other Related Sections (Optional)

Additional preprocessing techniques such as scaling or dimensionality reduction could be explored in future work to further improve model performance. For example, feature scaling may improve the performance of certain algorithms that are sensitive to the magnitude of input variables. Further analysis could also explore feature importance to identify which medical attributes contribute most significantly to thyroid disease prediction.

Chapter 3

Methodology

This section describes the machine learning framework used in this project, including the programming environment, data preparation steps, model selection, training process, and evaluation methods.

3.1 Introduction to Python for Machine Learning

Python is widely used for machine learning because of its extensive ecosystem of scientific computing libraries. In this project, Python was used as the primary programming language to implement the data preprocessing steps, machine learning models, and performance evaluation.

Several popular Python libraries were used, including:

- **Pandas** for data manipulation and preprocessing
- **NumPy** for numerical computations
- **Scikit-learn** for implementing machine learning models and evaluation metrics

These libraries provide efficient tools for building machine learning pipelines and performing model comparisons.

3.2 Platform and Machine Configurations Used

All experiments for this project were conducted on a personal computer using Python. The development environment included Visual Studio Code and a Python virtual environment configured using Conda.

The key software configuration used in the project includes:

- Python 3.x
- Pandas
- NumPy
- Scikit-learn

These tools provided the necessary functionality to preprocess the dataset, train machine learning models, and evaluate predictive performance.

3.3 Data Split

After preprocessing the dataset, the data was divided into training and testing sets. Data splitting is an essential step in machine learning because it allows the model to be evaluated on unseen data.

In this project, the dataset was split using an **80/20 train-test split**, where:

- **80% of the data** was used to train the machine learning models
- **20% of the data** was used to evaluate model performance

This approach ensures that the models are evaluated on data that was not used during training, providing a more realistic measure of predictive performance.

3.4 Model Planning

Three supervised machine learning classification models were selected for this project:

- Logistic Regression
- Decision Tree
- Random Forest

These models were chosen because they represent different machine learning approaches. Logistic Regression is a linear classification model, Decision Trees capture non-linear decision boundaries, and Random Forest is an ensemble method that combines multiple decision trees to improve prediction performance.

By comparing these models, the project aims to determine which approach performs best for predicting thyroid disease.

3.5 Model Training

Each model was trained using the training dataset after preprocessing was completed. The training process involves allowing the model to learn relationships between input features and the target variable.

The Scikit-learn library was used to implement the models. During training, the models adjusted their internal parameters in order to minimize classification errors and improve predictive accuracy.

Once the training process was complete, the models were applied to the testing dataset to generate predictions.

3.6 Model Evaluation

To evaluate the performance of each model, several common classification metrics were used. These include:

- **Accuracy**, which measures the proportion of correct predictions
- **Precision**, which measures how many predicted positive cases were correct
- **Recall**, which measures how well the model identifies actual positive cases
- **F1-score**, which provides a balance between precision and recall

These metrics provide a comprehensive evaluation of model performance, particularly in cases where the dataset may contain class imbalance.

3.7 Model Optimization

Model optimization involves adjusting model parameters and configurations to improve predictive performance. In this project, model parameters were configured using the default implementations provided by the Scikit-learn library.

Although extensive hyperparameter tuning was not performed, the comparison of multiple machine learning models provides insight into which algorithm is best suited for the thyroid disease prediction task.

3.8 Final Model Building

After training and evaluating the models, the final step was to compare the performance of each algorithm. The results showed that all models achieved high classification accuracy, with the Random Forest model achieving the best overall performance.

The final model comparison demonstrates how different machine learning techniques can be applied to medical datasets to identify patterns and predict disease outcomes.

Chapter 4

Results

This section presents the performance results of the machine learning models used in this project. The models were evaluated using the testing dataset to determine their ability to predict thyroid disease.

4.1 Description of the Models

Three supervised machine learning classification models were implemented and evaluated in this project:

- **Logistic Regression**: A statistical classification model that predicts the probability of a binary outcome using a logistic function. It is commonly used for medical classification problems because of its interpretability and efficiency.
- **Decision Tree**: A non-linear classification model that splits the dataset into branches based on feature values. Decision trees are capable of capturing complex relationships between variables and are easy to visualize and interpret.
- **Random Forest**: An ensemble learning method that builds multiple decision trees and combines their predictions. By aggregating the predictions of several trees, Random Forest models often achieve better accuracy and improved generalization compared to a single decision tree.

4.2 Performance Metrics

Several evaluation metrics were used to assess the predictive performance of the models:

- **Accuracy** measures the proportion of correct predictions made by the model.
- **Precision** measures the proportion of predicted positive cases that were actually correct.
- **Recall** measures how well the model identifies actual positive cases.
- **F1-score** provides a balance between precision and recall.

These metrics provide a more complete evaluation of model performance, particularly in datasets where class imbalance may exist.

4.3 Results Table

The performance of the evaluated machine learning models was measured using accuracy, precision, recall, and F1-score. These metrics provide a comprehensive evaluation of how well each model predicts thyroid disease cases.

Table 4.1: Performance comparison of machine learning models for thyroid disease detection.

Model	Accuracy	Precision	Recall	F1 Score
Logistic Regression	0.9839	0.875	0.467	0.609
Decision Tree	0.9839	0.714	0.667	0.690
Random Forest	0.9875	0.833	0.667	0.741

Table Table 4.1 summarizes the performance of the evaluated models. All three models achieved high accuracy, indicating strong predictive performance on the dataset. Among them, the Random Forest model achieved the highest overall performance, particularly in terms of F1-score, which reflects a better balance between precision and recall. This suggests that Random Forest is the most suitable model for thyroid disease detection in this study.

4.4 Interpretation of the Results

The results show that all three models achieved high classification accuracy, indicating that the dataset contains strong patterns that allow effective prediction of thyroid disease.

Among the evaluated models, the **Random Forest model achieved the highest accuracy and F1-score**, making it the best-performing model in this study. The ensemble nature of Random Forest allows it to combine multiple decision trees, which improves predictive performance and reduces overfitting.

The Decision Tree model also performed well and achieved higher recall than Logistic Regression. This suggests that the Decision Tree was more effective at identifying positive thyroid disease cases.

Logistic Regression achieved high overall accuracy but lower recall compared to the other models. This may indicate that the model was more conservative in predicting positive cases, potentially due to class imbalance in the dataset.

Overall, the results demonstrate that ensemble learning methods such as Random Forest can provide strong predictive performance when applied to medical classification problems.

4.5 Visualization

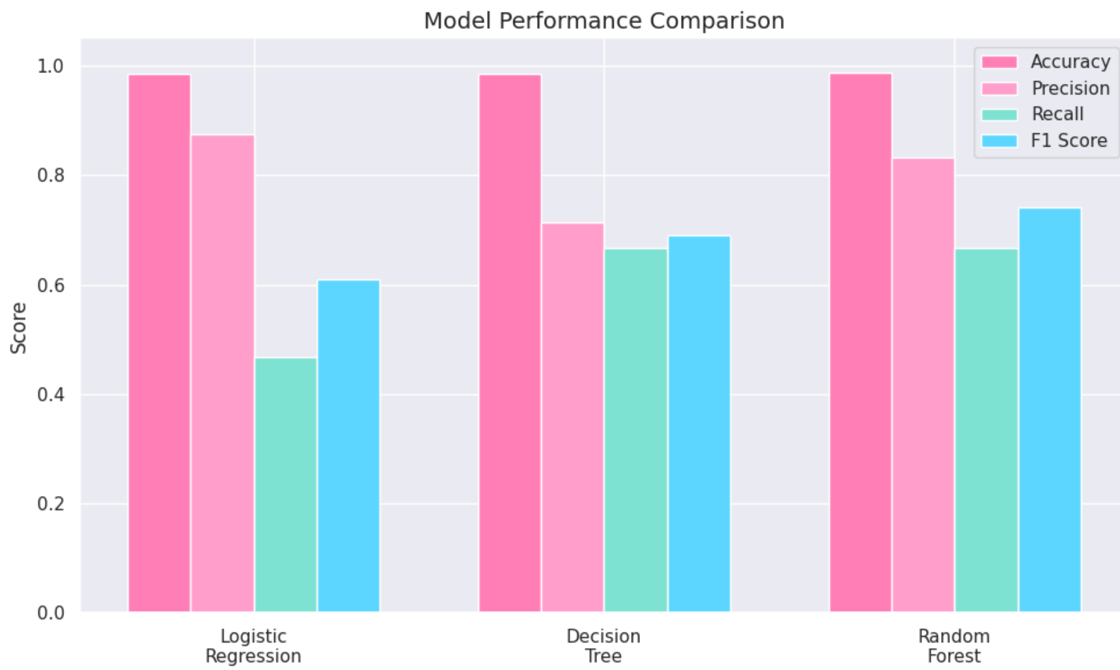

Figure 4.1: Performance comparison

Figure 1: Performance comparison of Logistic Regression, Decision Tree, and Random Forest models based on accuracy, precision, recall, and F1-score.

Figure 1 compares the performance of the three machine learning models used in this study. The Random Forest model achieved the highest overall performance with the best accuracy and F1-score among the evaluated models. Logistic Regression and Decision Tree also performed well, but Random Forest provided a better balance between precision and recall, making it the most suitable model for thyroid disease detection in this dataset.

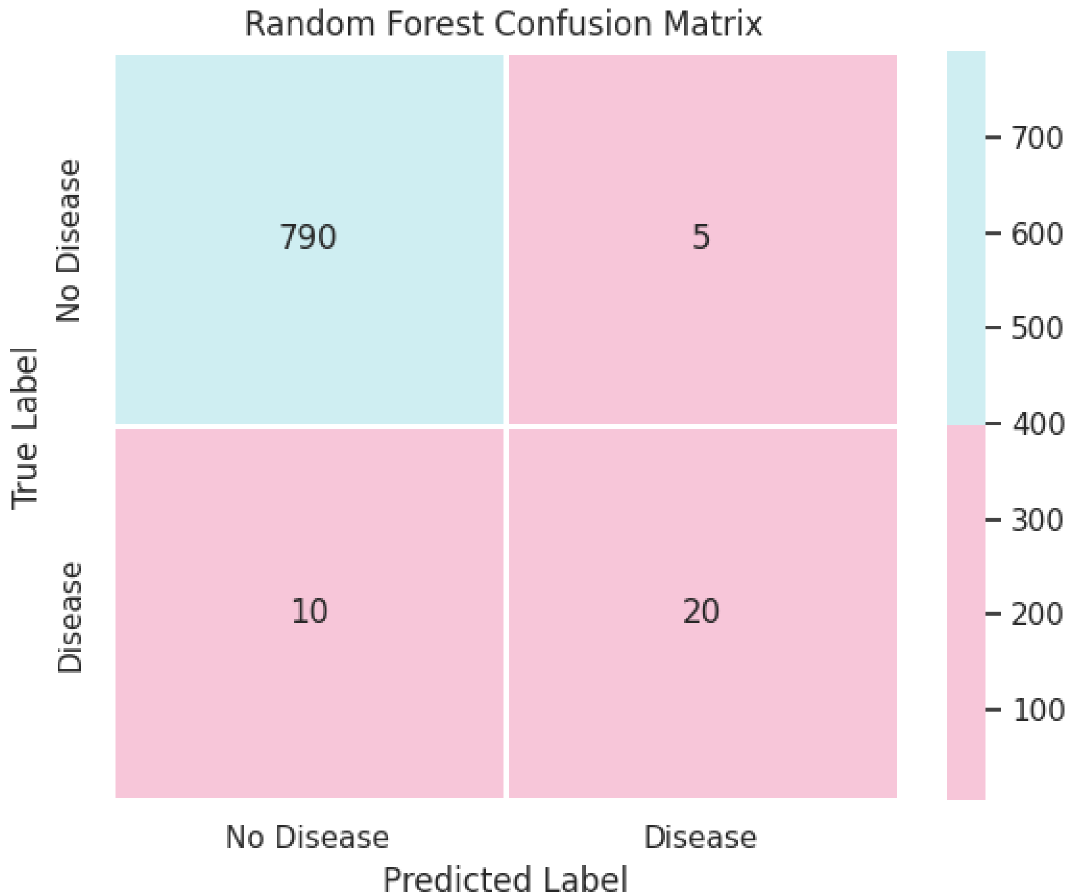

Figure 4.2: Confusion matrix

Figure 2: Confusion matrix for the Random Forest classifier showing predicted versus actual thyroid disease classifications.

Figure 2 illustrates the classification performance of the Random Forest model using a confusion matrix. The matrix shows that the model correctly classified the majority of non-disease cases while also identifying a portion of disease cases. However, some disease cases were misclassified, which may be due to class imbalance in the dataset. Despite this, the model demonstrates strong overall predictive performance.

4.6 Sensitivity Analysis

Sensitivity analysis was not extensively performed in this study. However, future work could explore hyperparameter tuning or feature selection techniques to further improve model performance and robustness.

Chapter 5

Conclusion

This project explored the use of machine learning techniques for predicting thyroid disease using the Thyroid Disease Dataset from the UCI Machine Learning Repository. The objective of the study was to evaluate whether patient medical attributes could be used to accurately classify hyperthyroid conditions.

The dataset was first preprocessed by handling missing values, converting variables to appropriate formats, and encoding categorical features. After preprocessing, the data was divided into training and testing sets in order to evaluate the performance of different machine learning models.

Three supervised learning algorithms were implemented and compared: Logistic Regression, Decision Tree, and Random Forest. These models were trained using the processed dataset and evaluated using several performance metrics, including accuracy, precision, recall, and F1-score.

The experimental results showed that all models achieved high predictive accuracy. Among the models tested, the Random Forest classifier achieved the best overall performance. The ensemble nature of Random Forest allows it to capture complex relationships in the data and produce more robust predictions compared to individual models.

The results of this study demonstrate that machine learning techniques can be effective tools for analyzing medical datasets and predicting thyroid disease conditions. Such models may assist healthcare professionals by providing decision-support tools that help identify potential disease cases.

Future work could explore additional improvements such as hyperparameter tuning, feature importance analysis, and addressing class imbalance in the dataset. Further experimentation with more advanced machine learning models or larger medical datasets could also improve predictive performance and provide deeper insights into thyroid disease detection.

References

UCI Machine Learning Repository. (n.d.). *Thyroid Disease Dataset*. https://archive.ics.uci.edu/ml/datasets/Thyroid+Disease

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., … & Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.